Woven Artificial Profile (WARP): Face Video Synthesis from Audio and Profile

By- B21SK05

DEMO

Demo of the fully functioning application.

Web application is created using FastAPI and React JS.

Link to Sample Output 1:
https://drive.google.com/file/d/1QUofEJGhsWNXOKQJqPo7k6zstDtvkwJO/view?usp=sharing

1 Introduction

- Motivation
- Problem Statement
- Demographic
- Literature Work

Motivation

Indirect vs. Direct Communication

- Indirect modes of communication lack the essence of Direct means- tone, expressions, gestures and intent.
- Information can be easily misinterpreted via indirect communication.
- Video/Audio calls, though a great alternative, are highly demanding of time and availability of all parties involved.
- WARP aims to personalise indirect communication.

Problem Statement-

Given just a face image and text, WARP is a highly versatile model that can generate realistic talking videos with audio.

Demographic

- Entertainment & Broadcasting Industries
- Educational Institutes
- Marketing Sectors
- Virtual Assistants & Avatars
- Remote Officials
- General Public

Literature Work

Neural Voice Puppetry [5]

Audio driven network that takes audio as input and uses an intermediate 3D face model that reconstructs the target person speaking and checks for synchronisation.

Disadvantage- Audio centric.

Temporal GAN [6]

Generative model capable of learning semantic representations unlabelled videos. Generates videos with number of frames equal to number of latent variables.

Disadvantage- Requires a large latent space. Audio inclusion is out of scope.

DualLip [7]

Leverages task duality by lip reading and generation using unlabelled text and lip videos. Synchronizes the lip generation model to the text.

Disadvantage- Replaces the lip on the guide face by vector concatenation.

2

The Product

Modules, Implementation & Results

Modules

- **Text-to-Speech**: *audio generation from text*
- **StyleGAN**: *fake face generation*
- **LipGAN**: *lip sync model*

WARP- Ensemble of the above three modules.

Text-to-Speech

- Converts a given text to audio.
- VoiceRSS API service.
- Cloud based API
- Offers several options for language and voice

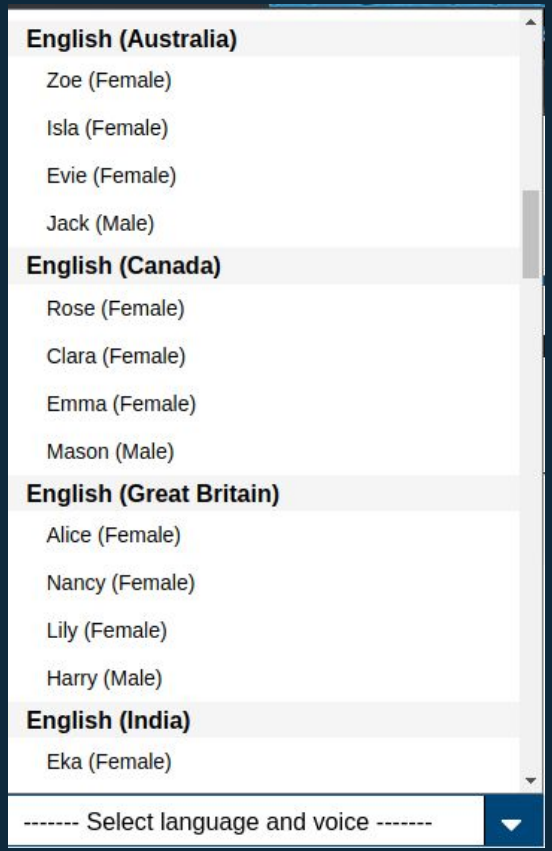

http://www.voicerss.org/

StyleGAN [1]

- A style-based generator architecture for Generative Adversarial Networks .
- Generates fake faces.
- Produces high quality Images of size 1024x1024.
- StyleGAN[1] is an improved version of ProGAN [4].
- Generator has 2 networks (mapping and synthesis).
- Adaptive Instance Normalization (AdaIN).
- Discriminator is similar to Synthesis network.
- Celeba HQ dataset.
- FID score of the GAN is 167.26 for a sample of 200 fake images.

Image credits [1]

LipGAN: Description

- [2]
- Generates lip synced transitional videos given an audio and source-destination images.
- Dataset- LRS2[3] dataset consisting of very short 1-10 second annotated videos.
- Generator: Audio Encoder, Face Encoder. For a given frame, output is a modified frame to sync with the corresponding audio segment.
- Discriminator: Audio and Face Encoder. Binary class output describing the level of synchronization.

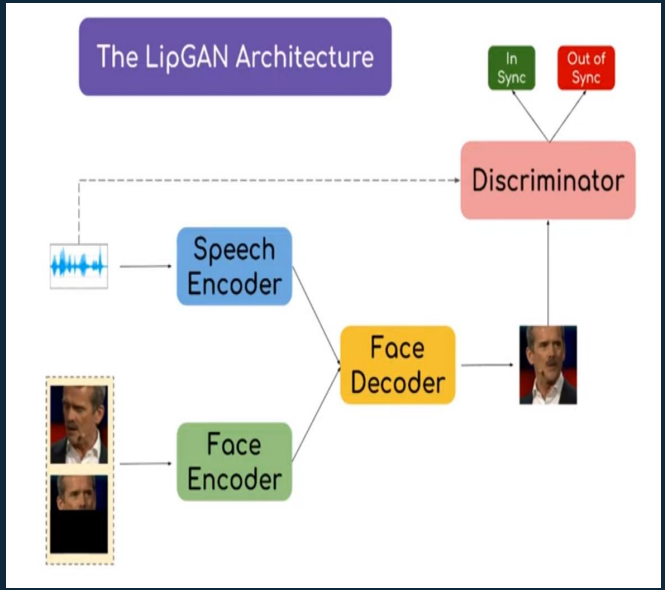

Image Source- [3]

LipGAN: Preprocessing for Training

- Original data is just simple videos.
- `Valid` frames: frames containing face views only.
- Use ***dlib***'s Frontal Face Detector (FFD) to extract *valid* frames and keep the ones with the maximum bounding box area of face view.
- Extract audio using ***ffmpeg***.
- Generate the Mel Spectrogram* and save the Mel Frequency Cepstral Coefficients (MFCC) features using ***Librosa***.

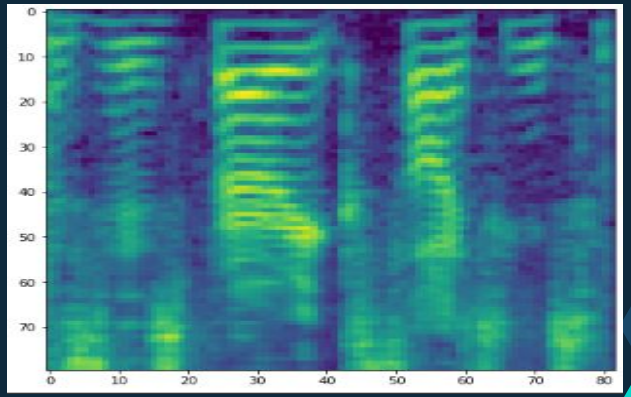

* Spectrogram** frequencies transformed to Mel Scale***.
** FFTs of audio signals stacked on top of each other.
*** scale of human audio perception rather than audio itself.

LipGAN: Training

- The model takes a random masked frame (other than the ground truth), target audio segment and ground truth frame as input.
- The training objective is to recreate a frame matching the ground truth frame for a given audio segment. Similar to Conditional GAN.
- A joint embedding is formulated by concatenating the face and audio encoder outputs.
- The decoder has a Sigmoid head for classifying the recreated frame as real/fake wrt. the ground truth.

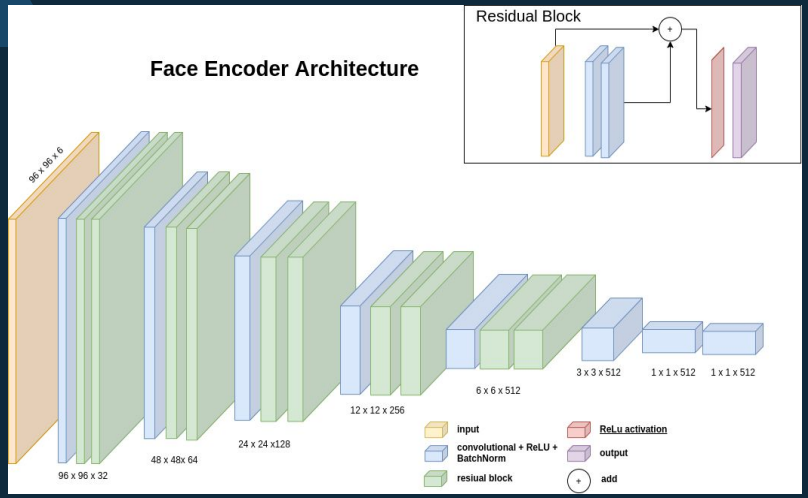

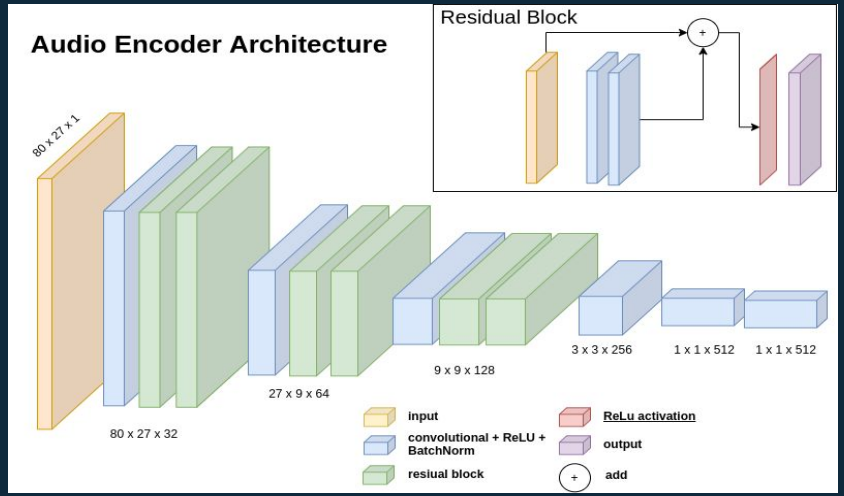

Architecture Diagrams for Face and Audio Encoders generated using https://app.diagrams.net/

WARP: Video Generation

- The model gets the text as input and generates the audio using TTS technique.
- If no user image is given, it generates a fake face for the video.
- Extracts the Mel Spectrogram from the audio and initializes the frames to match the audio length.
- For each frame, the model predicts the face view with accurate lip pose corresponding to the audio segment.
- Stitches and writes the video together using OpenCV.
- Uses ffmpeg to combine the audio and video for the final output.

3

Extras

- Timeline
- Contributions
- References

Contributions

Nethra Gunti

S20180010061

- LipGAN Preprocessing
- LipGAN Training
- Video Generation

Tarun Teja Obinna

S20180010120

- StyleGAN
- LipGAN Training
- Web Application

Nithin Ramancha

S20170010120

- Text-to-Speech
- Web App UI
- Web Application

References

1. Karras, T., Laine, S., & Aila, T. (2019). A style-based generator architecture for generative adversarial networks. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (pp. 4401-4410).
2. Prajwal K R, Rudrabha Mukhopadhyay, Jerin Philip, Abhishek Jha, Vinay Namboodiri, and C V Jawahar. 2019. Towards Automatic Face-to-Face Translation. In proceedings of the 27th ACM International Conference on Multimedia (MM '19).
3. T. Afouras, J. S. Chung, A. Senior, O. Vinyals, A. Zisserman, Deep Audio-Visual Speech Recognition, arXiv:1809.02108.
4. Karras, T., Aila, T., Laine, S., & Lehtinen, J. (2017). Progressive growing of gans for improved quality, stability, and variation. arXiv preprint arXiv:1710.10196.
5. Thies, J., Elgharib, M., Tewari, A., Theobalt, C., & Nießner, M. (2020, August). Neural voice puppetry: Audio-driven facial reenactment. In European Conference on Computer Vision (pp. 716-731). Springer, Cham.
6. Saito, M., Matsumoto, E., & Saito, S. (2017). Temporal generative adversarial nets with singular value clipping. In Proceedings of the IEEE international conference on computer vision (pp. 2830-2839).
7. Chen, W., Tan, X., Xia, Y., Qin, T., Wang, Y., & Liu, T. Y. (2020, October). DualLip: A System for Joint Lip Reading and Generation. In Proceedings of the 28th ACM International Conference on Multimedia (pp. 1985-1993).

Thanks!

Nethra Gunti (S20180010061)
Tarun Teja Obbina (S20180010120)
Nithin Ramancha (S20170010120)

Credits

Special thanks to all the people who made and released these awesome resources for free:

- Presentation template by SlidesCarnival
- Photographs by Unsplash